



VIT[®]
B H O P A L
www.vitbhopal.ac.in

VIT Bhopal University,
Bhopal-Indore Highway
Kothrikalan, Sehore, Madhya Pradesh - 466114

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
ACADEMIC YEAR: 2024-2025

CSE3009- PARALLEL AND DISTRIBUTED COMPUTING
LABORATORY RECORD

WINTER SEMESTER 2024 - 2025

Submitted By

Name : Saransh Prajapati

Reg No : 22BCE10795

Submitted to

Dr. Komarasamy G,
Senior Associate Professor, SCAI
VIT Bhopal University

DECLARATION

I **Saransh Prajapati (22BCE10795)** hereby declare that this record of observation is based on the experiments carried out and recorded by me during the laboratory Course of **CSE3009-PARALLEL AND DISTRIBUTED COMPUTING LABORATORY, VIT Bhopal University**, Kothrikalan, Sehore, Madhya Pradesh – 466114.

Date: _____

Signature of the student

Name of the Student: Saransh Prajapati

Reg Number : 22BCE10795

Countersigned by Staff

INDEX

E.NO	DATE	EXPERIMENT NAME	PG.NO	FACULTY SIGN
1	27/01/2025	OPENMP – BASIC PROGRAMS SUCH AS 1. VECTOR ADDITION, 2. DOT PRODUCT	1	
2	27/01/2025	OPENMP – 1. LOOP WORK-SHARING AND 2. SECTIONS WORK-SHARING		
3		OPENMP – 1. COMBINED PARALLEL LOOP REDUCTION AND 2. ORPHANED PARALLEL LOOP REDUCTION		
4		OPENMP – MATRIX MULTIPLY (SPECIFY RUN OF A GPU CARD, LARGE SCALE DATA ... COMPLEXITY OF THE PROBLEM NEED TO BE SPECIFIED)		
5		MPI – BASICS OF MPI		
6		MPI – COMMUNICATION BETWEEN MPI PROCESS		
7		MPI – COLLECTIVE OPERATION WITH "SYNCHRONIZATION"		
8		MPI – COLLECTIVE OPERATION WITH "DATA MOVEMENT"		
9		MPI – COLLECTIVE OPERATION WITH "COLLECTIVE COMPUTATION"		
10		MPI – NON-BLOCKING OPERATION		

Ex. No: 1	OPENMP – BASIC PROGRAMS SUCH AS 1. DOT PRODUCT, 2. VECTOR ADDITION.
Date:	

AIM: OpenMP – Basic programs such as Vector addition, Dot Product

ALGORITHM:

1. Dot Product:

- Input Vectors: Obtain two vectors a and b of equal length n .
- Initialize Result
- Iterate Over Elements:
 - For each index i from 0 to $n-1$:
 - Multiply $a[i]$ by $b[i]$ and add the result to result.
- Output Result.

2. Vector Addition:

- Define Constants: Set the size of the arrays and the number of threads.
- Allocate Memory: Dynamically allocate memory for the input arrays and the result array.
- Initialize Arrays with values.
- Parallelize Addition:
 - Use `#pragma omp parallel` for to distribute the addition operation across multiple threads.
- Free the dynamically allocated memory.

PROGRAM:

DOT-PRODUCT

```
#include <omp.h>
#include <stdio.h>

#define SIZE 100
#define CHUNK_SIZE 10

int main() {

    int index, total_elements;
    float vector_a[SIZE], vector_b[SIZE], dot_product = 0.0;
    total_elements = SIZE;

    for (index = 0; index < total_elements; index++) {
        vector_a[index] = index * 1.0f;
        vector_b[index] = index * 2.0f;
    }

    #pragma omp parallel for default(none) shared(vector_a, vector_b, total_elements) private(index)
    schedule(static, CHUNK_SIZE) reduction(+:dot_product)

    for (index = 0; index < total_elements; index++) {
        dot_product += (vector_a[index] * vector_b[index]);
    }
    printf("Dot Product Result: %f\n", dot_product);
}
```

VECTOR-ADDITION

```
#include <omp.h>
#include <stdio.h>

int main() {
    int vector_a[5] = {1, 4, 6, 7, 8};
    int vector_b[5] = {2, 3, 5, 9, 0};
    int vector_c[5] = {0, 0, 0, 0, 0};

    omp_set_num_threads(3);

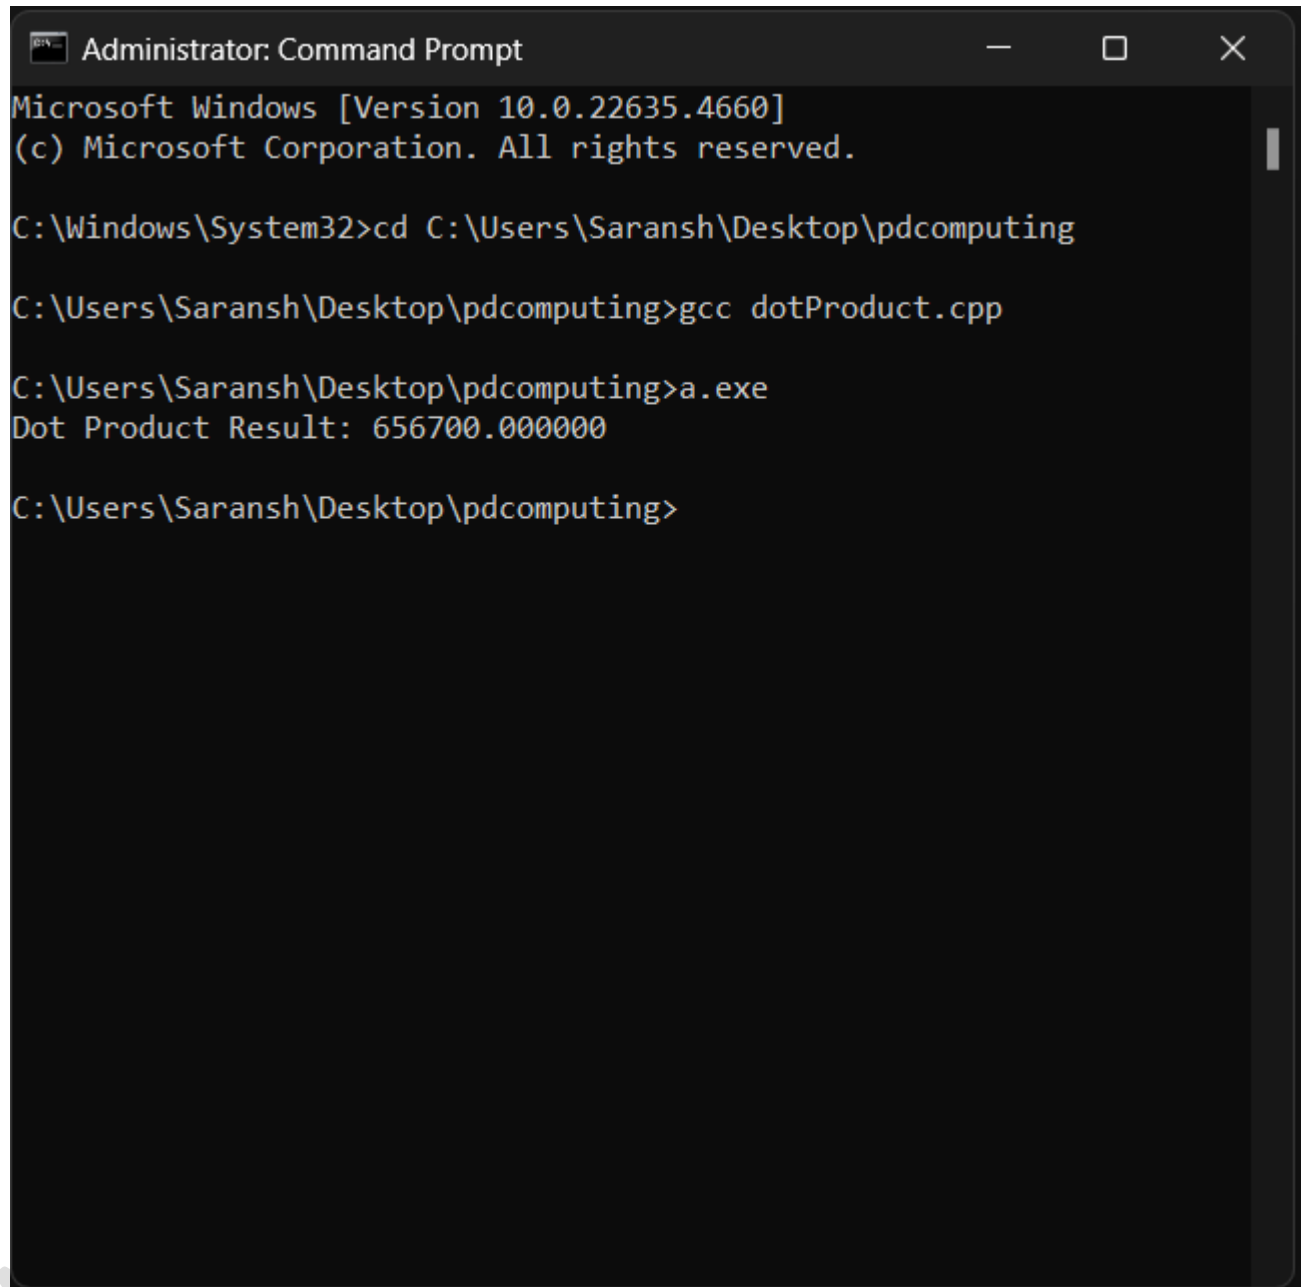
    #pragma omp parallel for shared(vector_a, vector_b, vector_c) schedule(static, 1)
    for (int idx = 0; idx < 5; idx++) {
        printf("Thread %d works on element %d\n", omp_get_thread_num(), idx);
        vector_c[idx] = vector_a[idx] + vector_b[idx];
    }

    printf("\n");
    printf("i \t vector_a[i] \t + \t vector_b[i] \t = \t vector_c[i] \n");
    for (int idx = 0; idx < 5; idx++) {
        printf("%d \t %d \t \t %d \t \t %d \n", idx, vector_a[idx], vector_b[idx], vector_c[idx]);
    }

    return 0;
}
```

OUTPUT:

1. Dot Product Output:



```
Administrator: Command Prompt
Microsoft Windows [Version 10.0.22635.4660]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>cd C:\Users\Saransh\Desktop\pdcomputing

C:\Users\Saransh\Desktop\pdcomputing>gcc dotProduct.cpp

C:\Users\Saransh\Desktop\pdcomputing>a.exe
Dot Product Result: 656700.000000

C:\Users\Saransh\Desktop\pdcomputing>
```

2. Vector Addition Output:

```
Administrator: Command Prompt
Microsoft Windows [Version 10.0.22635.4660]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>cd C:\Users\Saransh\Desktop\pdcomputing

C:\Users\Saransh\Desktop\pdcomputing>gcc dotProduct.cpp

C:\Users\Saransh\Desktop\pdcomputing>a.exe
Dot Product Result: 656700.000000

C:\Users\Saransh\Desktop\pdcomputing>gcc -fopenmp vectorAdd.cpp

C:\Users\Saransh\Desktop\pdcomputing>a.exe
Thread 0 works on element 0
Thread 0 works on element 3
Thread 2 works on element 2
Thread 1 works on element 1
Thread 1 works on element 4

i      vector_a[i]  +      vector_b[i]  =      vector_c[i]
0      1            2            3
1      4            3            7
2      6            5            11
3      7            9            16
4      8            0            8

C:\Users\Saransh\Desktop\pdcomputing>_
```


RESULT:

1. Dot Product Result:

The dot product of the two vectors vector A and vector B is 656700, indicating the sum of the element-wise products of the two vectors.

2. Vector Addition Result:

The resulting vector C is obtained by adding each element of vector A to its corresponding element in vector B.

Ex. No: 2	OpenMP – Loop work-sharing and sections work-sharing
Date: 27/1/25	

AIM: OpenMP – Loop work-sharing and sections work-sharing

ALGORITHM:

1. Loop work-sharing:

- Uses #pragma omp parallel to create a parallel region
- Distribute iterations: Assign a starting and ending iteration to each thread. This can be done statically (pre-determined ranges) or dynamically (work is assigned as threads finish).
- Parallel Execution: Each thread executes the loop iterations assigned to it.

2. Sections Work-Sharing:

- Assign Sections: Each section of code is assigned to a separate thread.
- Parallel Execution: Each thread executes the section assigned to it.
- Implicit Barrier: All threads wait until all sections have completed execution.
- Continue execution.

PROGRAM:

Loop Work-Sharing

```
#include <iostream>
#include <omp.h>
#include <stdio.h>

int main() {
    int n = 10;
    int arr[n];
    #pragma omp parallel for
    for (int i = 0; i < n; i++) {
        arr[i] = i * 2;
        cout << "Thread " << omp_get_thread_num() << " processed element " << i << endl;
    }
    return 0;
}
```

Sections Work-Sharing

```
#include <iostream>
#include <stdio.h>
#include <omp.h>

void function1() {
    cout << "Thread " << omp_get_thread_num() << " executing function1" << endl;
}

void function2() {
    cout << "Thread " << omp_get_thread_num() << " executing function2" << endl;
}

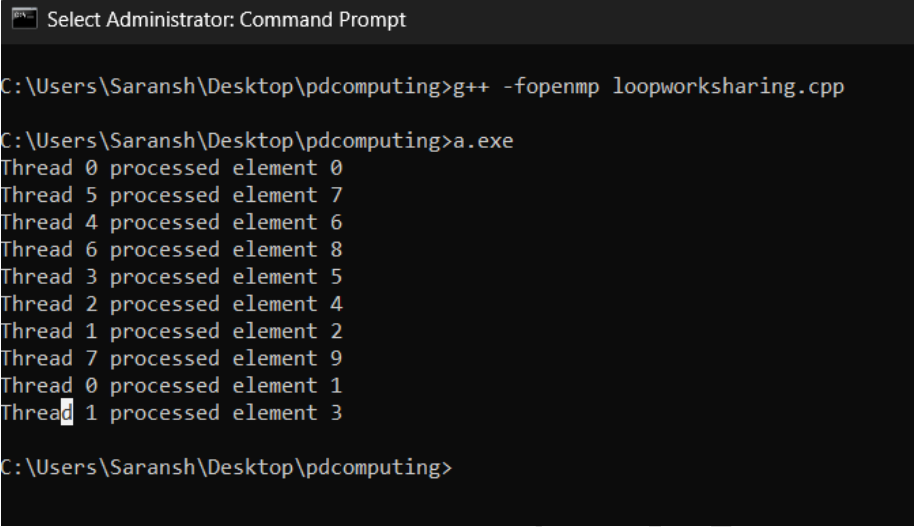
int main() {
    #pragma omp parallel sections
    {
        #pragma omp section
        function1();

        #pragma omp section
        function2();
    }

    return 0;
}
```

OUTPUT:

Loop Work-Sharing:



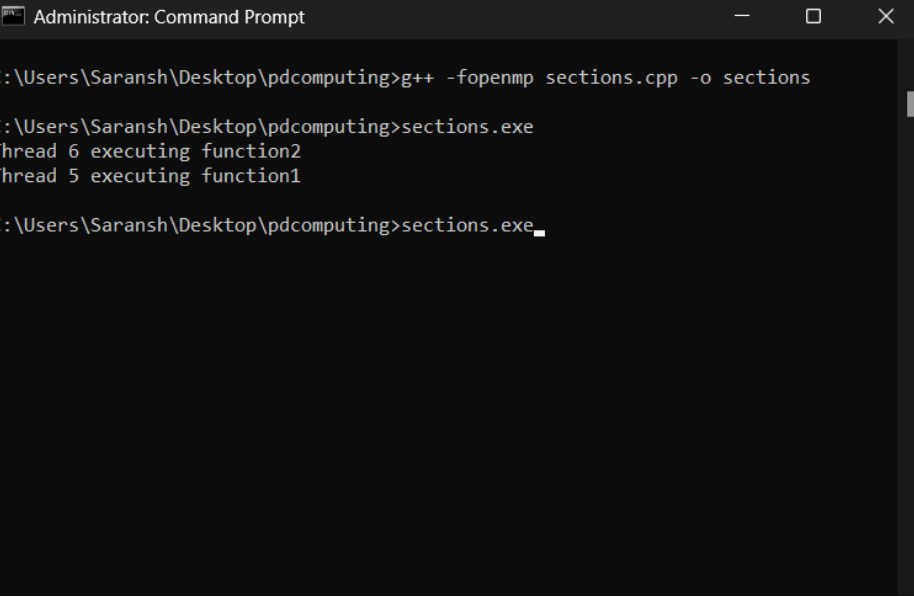
```
Select Administrator: Command Prompt

C:\Users\Saransh\Desktop\pdcomputing>g++ -fopenmp loopworksharing.cpp

C:\Users\Saransh\Desktop\pdcomputing>a.exe
Thread 0 processed element 0
Thread 5 processed element 7
Thread 4 processed element 6
Thread 6 processed element 8
Thread 3 processed element 5
Thread 2 processed element 4
Thread 1 processed element 2
Thread 7 processed element 9
Thread 0 processed element 1
Thread 1 processed element 3

C:\Users\Saransh\Desktop\pdcomputing>
```

Sections Work-Sharing:



```
Administrator: Command Prompt

C:\Users\Saransh\Desktop\pdcomputing>g++ -fopenmp sections.cpp -o sections

C:\Users\Saransh\Desktop\pdcomputing>sections.exe
Thread 6 executing function2
Thread 5 executing function1

C:\Users\Saransh\Desktop\pdcomputing>sections.exe
```

RESULT:

Both Loop work-sharing & sections work sharing are successfully implemented.

Ex. No: 3	OpenMP – Combined parallel loop reduction and Orphaned parallel loop reduction
Date: 12/2/25	

AIM: OpenMP – Combined parallel loop reduction and Orphaned parallel loop reduction.

ALGORITHM:

1. Combined parallel loop reduction:

- The sum variable is initialized before the parallel region. This is important.
- `#pragma omp parallel for reduction(+: sum)`: This directive tells OpenMP to parallelize the loop and perform a reduction on the sum variable using the + operator (addition).
- Inside the loop, each thread adds its portion of the array to its private copy of sum.
- After the loop, OpenMP automatically adds all the private sum values together and stores the final result in the original sum variable

2. Orphaned parallel loop reduction:

- Assign Sections: Each section of code is assigned to a separate thread.
- Data Initialization: Create and initialize the data that will be processed in parallel.
- Parallel Region (Orphaned): Start a parallel region using `#pragma omp parallel`. The loop that does the parallel work will be *outside* this region (hence "orphaned").
- Parallel Loop (`#pragma omp for`): Use `#pragma omp for` *outside* the parallel region to distribute the loop iterations among the threads. This is the "orphaned" part.
- Thread-Local Variables: Inside the parallel region, declare any variables that need to be private to each thread. This is where you'd perform thread-specific setup.
- Computation (Inside the Loop): The actual computation happens within the loop, which is being executed in parallel. Each thread works on its assigned chunk of the loop.

- Reduction (If Needed): If you need to combine results from each thread (like summing up values), use a reduction clause or a critical section (although reduction is generally preferred for performance).
- Synchronization (If Needed): If threads need to coordinate or wait for each other, use OpenMP synchronization constructs (e.g., barriers).
- Result Collection (If Needed): Gather the results from each thread after the parallel loop has finished.
- Cleanup: Perform any necessary cleanup.

PROGRAM:

Combined parallel loop reduction:

```
#include <iostream>
#include <omp.h>
#include <stdio.h>
using namespace std;
```

```
int main() {
    int n = 1000;
    int arr[n];
    for (int i = 0; i < n; i++) {
        arr[i] = i + 1;
    }

    int sum = 0;
    #pragma omp parallel for reduction(+: sum)
    for (int i = 0; i < n; i++) {
        sum += arr[i];
    }

    cout << "Sum = " << sum << endl;
    return 0;
}
```


Orphaned parallel loop reduction:

```
#include <iostream>
#include <omp.h>
#include <vector>
#include <cmath>
using namespace std;

int main() {
    int n = 100000;
    vector<double> data(n);

    // 1. Data Initialization
    for (int i = 0; i < n; ++i) {
        data[i] = i * 0.01; // Example data
    }
    double sum_of_sin = 0.0;
    // 2. Parallel Region (Orphaned)
    #pragma omp parallel
    {
        // 4. Thread-Local Variables (if needed)
        double local_sum = 0.0;
        // 3. Parallel Loop (Orphaned)
        #pragma omp for
        for (int i = 0; i < n; ++i) {
            // 5. Computation (Inside the Loop)
            local_sum += sin(data[i]);
        }

        // 6. Reduction (using critical section for illustration - reduction is better!)
        #pragma omp critical
        {
            sum_of_sin += local_sum;
        }
    } // End of parallel

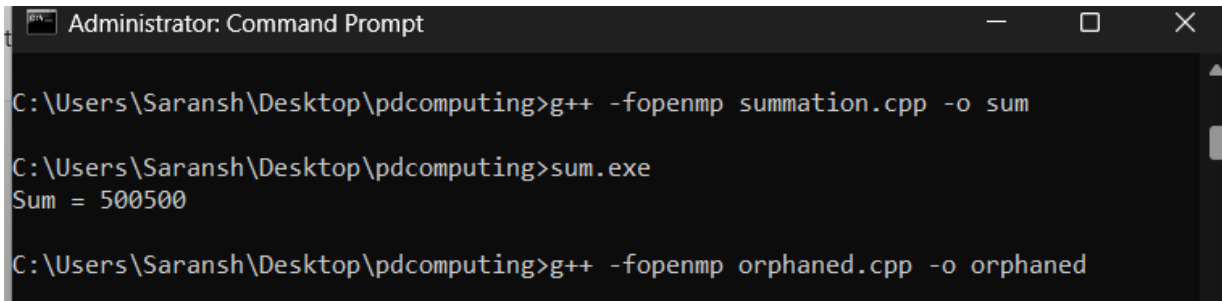
    // 8. Result Collection (already done with reduction)

    cout << "Sum of sin(data[i]): " << sum_of_sin << endl;

    return 0;
}
```

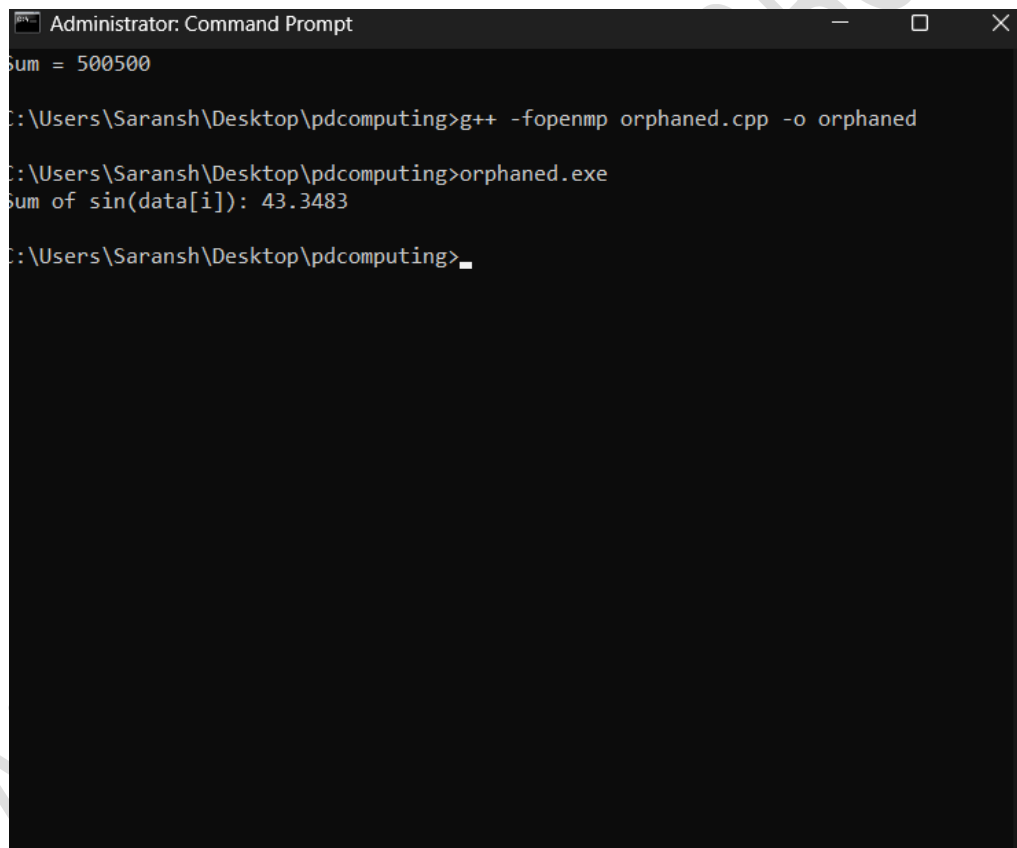
OUTPUT:

Combined parallel loop reduction:



```
Administrator: Command Prompt
C:\Users\Saransh\Desktop\pdcomputing>g++ -fopenmp summation.cpp -o sum
C:\Users\Saransh\Desktop\pdcomputing>sum.exe
Sum = 500500
C:\Users\Saransh\Desktop\pdcomputing>g++ -fopenmp orphaned.cpp -o orphaned
```

Orphaned parallel loop reduction:



```
Administrator: Command Prompt
Sum = 500500
C:\Users\Saransh\Desktop\pdcomputing>g++ -fopenmp orphaned.cpp -o orphaned
C:\Users\Saransh\Desktop\pdcomputing>orphaned.exe
Sum of sin(data[i]): 43.3483
C:\Users\Saransh\Desktop\pdcomputing>
```

RESULT:

Both Combined parallel loop reduction & Orphaned parallel loop reduction are successfully implemented.

Ex. No: 4	OpenMP – Matrix multiply (specify run of a GPU card, large scale data ... Complexity of the problem need to be specified)
Date: 12/2/25	

AIM: OpenMP – Matrix multiply (specify run of a GPU card, large scale data ... Complexity of the problem need to be specified)

ALGORITHM:

1. Combined parallel loop reduction:

- The sum variable is initialized before the parallel region. This is important.
- **Input:** Two matrices A (n x m) and B (m x p).
- **Output:** Result matrix C (n x p), where $C[i][j]$ is the dot product of the i-th row of A and the j-th column of B.
- **Initialization:** Create a matrix C of size n x p and initialize all its elements to 0.
- **Multiplication:**
 - Iterate through each row i of matrix A (from 0 to n-1).
 - Iterate through each column j of matrix B (from 0 to p-1).
 - Iterate through each element k in the i-th row of A and the k-th row of B (from 0 to m-1).
 - Calculate the dot product: $C[i][j] += A[i][k] * B[k][j]$.
- **Parallelization:** The outermost loop (iterating through rows of A) is parallelized using OpenMP's `#pragma omp parallel for` directive. This distributes the computation of different rows of C among multiple threads.
- **Return:** The resulting matrix C.

PROGRAM:

Orphaned parallel loop reduction:

```
#include <iostream>
#include <iostream>
#include <vector>
#include <omp.h>
using namespace std;

void matrixMultiply(const vector<vector<int>>& A, const vector<vector<int>>& B, vector<vector<int>>& C) {
    int n = A.size(); // Rows of A and C
    int m = B.size(); // Rows of B
    int p = B[0].size(); // Columns of B and C

    if (m != n) {
        cerr << "Matrices are not compatible for multiplication." << endl;
        return;
    }

    C.resize(n, vector<int>(p, 0));

    // Parallelize the outer loop (rows of A)

    #pragma omp parallel for
    for (int i = 0; i < n; ++i) {
        for (int j = 0; j < p; ++j) {
            for (int k = 0; k < m; ++k) {
                C[i][j] += A[i][k] * B[k][j];
            }
        }
    }
}

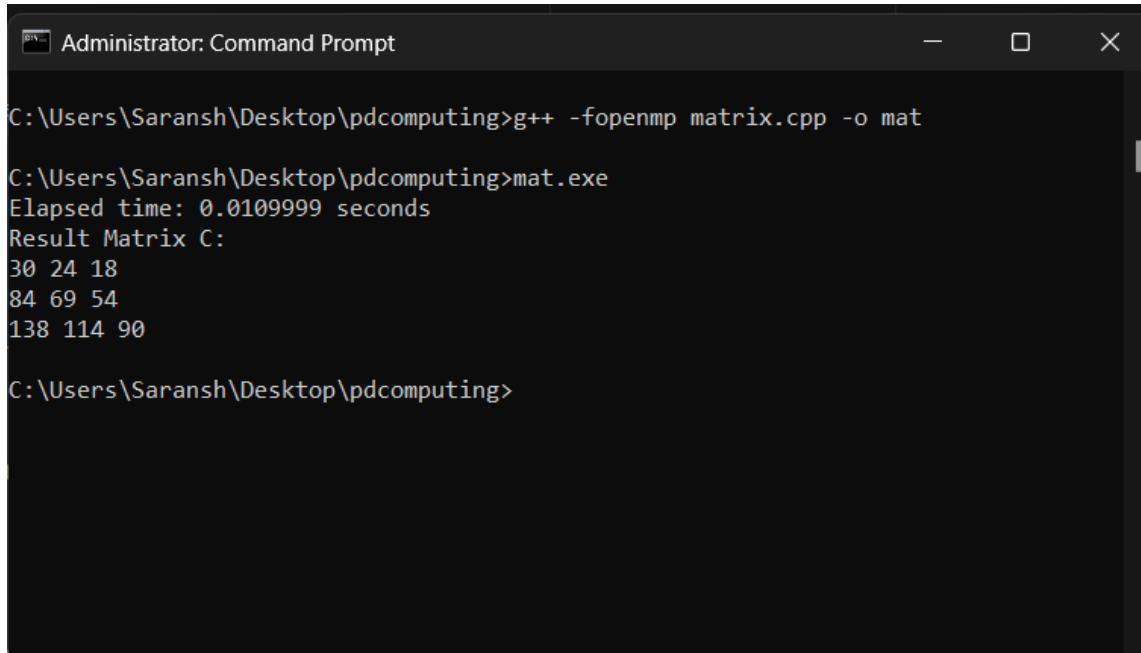
int main() {
    vector<vector<int>> A = {
        {1, 2, 3},
        {4, 5, 6},
        {7, 8, 9}
    };

    vector<vector<int>> B = {
        {9, 8, 7},
```

```
{6, 5, 4},
{3, 2, 1}
};
vector<vector<int>> C; // Result matrix
double start_time = omp_get_wtime();
matrixMultiply(A, B, C);
double end_time = omp_get_wtime();
double elapsed_time = end_time - start_time;
cout << "Elapsed time: " << elapsed_time << " seconds" << endl;
cout << "Result Matrix C:" << endl;

for (const auto& row : C) {
    for (int val : row) {
        cout << val << " ";
    }
    cout << endl;
}
return 0;
}
```

OUTPUT:



```
Administrator: Command Prompt

C:\Users\Saransh\Desktop\pdcomputing>g++ -fopenmp matrix.cpp -o mat

C:\Users\Saransh\Desktop\pdcomputing>mat.exe
Elapsed time: 0.0109999 seconds
Result Matrix C:
30 24 18
84 69 54
138 114 90

C:\Users\Saransh\Desktop\pdcomputing>
```

RESULT: